


SmolVLA inference latency: all devices on one scale



Shared linear scale; subsecond paths are visually small. Colors follow execution order, and labels show end-to-end medians only.