


SmolVLA inference latency by stage

RK3588 NPU denoising removes the largest bottleneck; multimodal prefix now dominates.



Stages follow execution order. Panels use different labeled linear scales; values are warmed medians. Other is the residual needed to match the end-to-end median.