



SmolVLA: share of inference time within each device

Slices follow execution order clockwise: vision, prefill, denoising, other.


RK3588 CPU




RK3588 NPU vision + CPU prefill + NPU denoising




RK3588 NPU vision + prefill + NPU denoising




i5-13490F CPU



MacBook Pro (Apple M5 Pro, 16-core GPU)



RTX 5060 CUDA



Warmed stage medians; percentages use end-to-end median as denominator. With RK3588 vision, prefill and denoising on NPU, vision is now the largest stage.